

# kaggle-v6.0

December 17, 2018

```
In [1]: from conf import *

In [2]: %load_ext autoreload
        %autoreload 2
        %reload_ext autoreload

In [3]: #eccezioni, non da dropare
        target = ["fullVisitorId", "totals.totalTransactionRevenue"]
        #num cols
        nums = ["visitStartTime", "totals.totalTransactionRevenue"]
        #cat cols
        cats = ["trafficSource.adwordsClickInfo.gclid", "trafficSource.referralPath", "trafficSource"]

        parameters = {
            #numero massimo di valori in una singola colonna per essere flattata, altrimenti d
            "max_new_feat": 500,
            #inviare a kaggle tramite l' API
            "commit": 0,
            #lgbm tuning parameters
            "n_leaves" : 512,
            "feature_fraction" : 0.99,
            "bagging_fraction" : 0.99,
            "learn_rate" : 0.004,
            #train rows
            "train_rows" : 10000,
            #test_rows, per submittare deve essere settato a -1
            "test_rows" : 500,
            #il metodo principale è lgbm ma si può testare anche la regressione lineare
            "test_also_lin_reg" : 1,
            #bagging frequency
            "bagging_freq" : 1,
            #transactionRevenue
            "transactionRevenue" : 0,
            #percentuale di dev e val
            "percentage" : 18,
            #grouping_mode_cats
            "grouping_mode_cats" : "mean",
```

```

        #score
        "final_score" : -1,
        #minio di alberi per il rf
        "min_child_samples" : -1
    }

    locals().update(parameters)

In [4]: if (commit==1) & (test_rows != -1):
        raise Exception("per submittare devi usare tutte le righe del test")

In [5]: %time train_df = load_df(train_file,train_rows,target)

Loaded train_v2.csv. Shape: (10000, 92)
CPU times: user 23.2 s, sys: 594 ms, total: 23.8 s
Wall time: 23.9 s

In [6]: %time test_df = load_df(test_file,test_rows,target)

Loaded test_v2.csv. Shape: (500, 76)
CPU times: user 2.72 s, sys: 47.6 ms, total: 2.77 s
Wall time: 2.78 s

In [7]: if "totals.totalTransactionRevenue" in test_df.columns:
        test_df = test_df.drop("totals.totalTransactionRevenue",axis=1)

In [8]: if not transactionRevenue:
        if "totals.transactionRevenue" in test_df.columns:
            test_df = test_df.drop("totals.transactionRevenue",axis=1)

In [9]: if not transactionRevenue:
        if "totals.transactionRevenue" in train_df.columns:
            train_df = train_df.drop("totals.transactionRevenue",axis=1)

In [10]: #controllare che nel nuovo test non droppi totalsRevenue perchè c'è solo nel train
train_df,test_df = drop_uncommons(train_df,test_df,target)

nums = find_num_cols(train_df,target,cats,nums)

cats = find_cat_cols(train_df,target,nums)

In [11]: cc(train_df,test_df,cats,nums)

In [12]: for col in nums:
        if (col not in target) | (col=="totals.totalTransactionRevenue"):
            train_df[col] = train_df[col].astype(float)
            if col!="totals.totalTransactionRevenue":
                test_df[col] = test_df[col].astype(float)

```

```

In [13]: # Impute 0 for missing target values
         train_df.fillna(0,inplace=True)
         test_df.fillna(0,inplace=True)

In [14]: train_df = stringify_cats(train_df,cats)
         test_df = stringify_cats(test_df,cats)

In [15]: #droppo le colonne che hanno troppa varianza
         train_df,test_df,cats = drop_exceeding(train_df,test_df,max_new_feat,cats,target)

In [16]: cc(train_df,test_df,cats,nums)

In [17]: train_df[cats] = train_df[cats].astype(str)
         test_df[cats] = test_df[cats].astype(str)

In [18]: #####QUI aggiungo il weekday

         #train_df["date"].weekday()

In [19]: %time train_df,test_df,cats = encode_cats(train_df,test_df,cats)

[*****]CPU times: user 4.67 s, sys: 2.04 s, total: 6.72 s
Wall time: 6.73 s

In [20]: cc(train_df,test_df,cats,nums)

In [21]: %time train_df = group_me(train_df,"fullVisitorId",cats,nums,grouping_mode_cats)

[*****]CPU times: user 455 ms, sys: 84.7 ms, total: 539 ms
Wall time: 537 ms

In [22]: %time test_df = group_me(test_df,"fullVisitorId",cats,nums,grouping_mode_cats)

[*****]CPU times: user 268 ms, sys: 21.4 ms, total: 289 ms
Wall time: 285 ms

In [23]: cc(train_df,test_df,cats,nums)

In [24]: load = -1
         base = "./saved_conf/"

         if load == 1:
             print("hai scelto di importare il dataset da disco")
             train_df = pd.read_csv(base + "dump_train")
             test_df = pd.read_csv(base + "dump_test")
             with open(base + "dump_parameters", 'r') as file:
                 file.read(json.loads(parameters))

```

```

        locals().update(parameters)
    else:
        if load == 0:
            print("hai scelto di scrivere il dataset su disco")
            %time train_df.to_csv(path_or_buf=base + "dump_train", header=True, mode='w',
                                test_df.to_csv(path_or_buf=base + "dump_test", header=True, mode='w', index=False),
                                with open(base + "dump_parameters", 'w') as file:
                                    file.write(json.dumps(parameters))
        else:
            print("hai scelto di non caricare nè scaricare il dataset")

```

hai scelto di non caricare nè scaricare il dataset

```

In [25]: train_id = train_df["fullVisitorId"].values
        test_id = test_df["fullVisitorId"].values

```

```

In [26]: #pulizia delle colonne con nomi assurdi
        #questa operazione può essere fatta in maniera safe perchè
        #a questo punto i due datasets hanno le stesse colonne con gli stessi nomi

```

```

In [27]: train_df.columns = [col[:30] for col in train_df.columns]
        test_df.columns = [col[:30] for col in test_df.columns]

```

```

In [28]: common_feats = list((set(train_df.columns).intersection(set(test_df.columns))).difference())

```

```

In [29]: #qui viene bloccato il controllo di coerenza poichè le colonne cambiano, in particolare
        #vengono accorciati ma è safe farlo perchè i nomi delle colonne sono importanti solo per la lettura

```

```

In [30]: #cc(train_df, test_df, cats, nums)

```

```

In [31]: #train_df.head()

```

```

In [32]: #splitto il dataframe in development e validation ma cercando di mantenere in maniera
        #corretta il rapporto dei compratori che è circa dell' 1%

```

```

#posso fare confronti con 0 perchè prima tutte le colonne sono state messe a 0 perciò

```

```

#divido il train in 2 parti: quelli che hanno speso che sono l' 1% e quelli che non hanno
        #speso 99% e da ognuno estraggo il tot% quindi mantengo il rapporto tra i due

```

```

train_money = train_df[train_df["totals.totalTransactionRevenue"]>0]
train_no_money = train_df[train_df["totals.totalTransactionRevenue"]==0]

```

```

percent = int(len(train_money)*percentage/100)
train_money_val = train_money.iloc[:percent,]

```

```

dev_df = train_money.iloc[percent:len(train_money),]
val_df = train_money_val

```

```

percent = int(len(train_no_money)*percentage/100)
train_no_money_val = train_no_money.iloc[:percent,]

dev_df = dev_df.append(train_no_money.iloc[percent:len(train_no_money),])
val_df = val_df.append(train_no_money_val )

#####
#voglio lavorare su un subset perciò provo a ridurre la grandezza
#mantengo il rapporto ma perdo info nelle features
#quantity=1

#dev_df=dev_df.iloc[:int(len(dev_df)*quantity),:]
#val_df=val_df.iloc[:int(len(val_df)*quantity),:]
#####

#dev_y contiene la colonna addestramento in dev già log1p
dev_y = np.log1p(dev_df["totals.totalTransactionRevenue"].values)
#val_y contiene la colonna target in val già log1p
val_y = np.log1p(val_df["totals.totalTransactionRevenue"].values)

#dev_x contiene colonne numeriche e cat senza transRev
dev_X = dev_df[ common_feats ]
#val_x contiene colonne numeriche e cat senza transRev
val_X = val_df[ common_feats ]
#test è ciò che dobbiamo trovare
test_X = test_df[ common_feats ]

```

```

In [33]: def write(tipo):
    parameters["final_score"] = final_score
    try:
        if len(pd.read_csv("./tests.csv").columns) != len(parameters.keys())+1:
            print("il file tests.csv contiene meno colonne del necessario, verrà sostituito")
            !rm "./tests.csv"
    except:
        print("il file tests.csv verrà creato ora perchè non esistente")
    with open("./tests.csv",'a') as ff:
        if os.fstat(ff.fileno()).st_size == 0:
            for k in parameters.keys():
                print(k+',',file = ff,sep='',end=' ' )
            print("type",file = ff)
            for v in parameters.values():
                print(str(v)+',',file = ff,sep='',end=' ' )
            print(tipo,file = ff )

```

## 0.1 Linear Regression

```
In [34]: from regression import lin
```

```
if test_also_lin_reg == 1:
    pred_test = lin(dev_X, dev_y, test_X)
    pred_val = lin(dev_X, dev_y, val_X)

    val_pred_df = pd.DataFrame({"fullVisitorId": val_df["fullVisitorId"].values})
    val_pred_df["totals.totalTransactionRevenue"] = val_df["totals.totalTransactionRevenue"]
    val_pred_df["PredictedRevenue"] = np.expml(pred_val)
    val_pred_df = val_pred_df.groupby("fullVisitorId")["totals.totalTransactionRevenue"]
    val_pred_df[val_pred_df["PredictedRevenue"] > 10^20] = 0

    final_score = np.sqrt(metrics.mean_squared_error(np.log1p(val_pred_df["totals.totalTransactionRevenue"]),
                                                         np.log1p(val_pred_df["PredictedRevenue"])))

    print(final_score)
    write("lin_reg")
```

```
1.0230158552172006
```

## 0.2 LightGBM single-tree

```
In [35]: # custom function to run light-gbm model
```

```
def lgbm(train_X, train_y, val_X, val_y, test_X):

    params = {
        "objective" : "regression",
        "metric" : "rmse",
        "num_leaves" : n_leaves,
        "feature_fraction" : feature_fraction,
        "bagging_fraction" : bagging_fraction,
        "bagging_freq" : bagging_freq,
        "learning_rate" : learn_rate,
        "verbosity" : -1
    }

    lgtrain = lgb.Dataset(train_X, label=train_y)
    lgval = lgb.Dataset(val_X, label=val_y)
    model = lgb.train(params, lgtrain, 10000, valid_sets=[lgval], early_stopping_rounds=100)

    pred_test_y = model.predict(test_X, num_iteration=model.best_iteration)
    pred_val_y = model.predict(val_X, num_iteration=model.best_iteration)
    return pred_test_y, model, pred_val_y
```

```
In [36]: #%time pred_test, model, pred_val = lgbm(dev_X, dev_y, val_X, val_y, test_X)
```

```
In [37]: def score():
    pred_val[pred_val < 0] = 0
```

```

val_pred_df = pd.DataFrame({"fullVisitorId":val_df["fullVisitorId"].values})
val_pred_df["totals.totalTransactionRevenue"] = val_df["totals.totalTransactionRevenue"]
val_pred_df["PredictedRevenue"] = np.expm1(pred_val)
val_pred_df = val_pred_df.groupby("fullVisitorId")["totals.totalTransactionRevenue"].transform(lambda x: np.exp(x))
final_score = np.sqrt(metrics.mean_squared_error(np.log1p(val_pred_df["totals.totalTransactionRevenue"]), val_df["totals.totalTransactionRevenue"]))
print(final_score)

sub_df = pd.DataFrame({"fullVisitorId":test_id})
pred_test[pred_test<0] = 0
sub_df["PredictedLogRevenue"] = np.expm1(pred_test)
sub_df = sub_df.groupby("fullVisitorId")["PredictedLogRevenue"].sum().reset_index()
sub_df.columns = ["fullVisitorId", "PredictedLogRevenue"]
sub_df["PredictedLogRevenue"] = np.log1p(sub_df["PredictedLogRevenue"])

write("LightGBM")

if commit:
    !kaggle competitions submit -c ga-customer-revenue-prediction -f {my_submission_file}

return final_score,sub_df

```

```

In [38]: def write_df(sub_df):
        sub_df.to_csv(path_or_buf=my_submission_file, header=True, mode='w',index=False)
        !wc -l {my_submission_file}

```

```

In [39]: #final_score,sub_df = score()

```

```

In [40]: #write_df(sub_df)

```

```

In [41]: def plot_imp(model):
        fig, ax = plt.subplots(figsize=(12,18))
        lgb.plot_importance(model, max_num_features=50, height=0.8, ax=ax)
        ax.grid(False)
        plt.title("LightGBM - Feature Importance", fontsize=15)
        plt.show()

```

```

In [42]: #plot_imp(model)

```

### 0.3 LightGBM con rf

```

In [43]: parameters["n_leaves"] = 400
        parameters["bagging_fraction"] = 0.99
        parameters["feature_fraction"] = 0.99
        parameters["bagging_freq"] = 20
        parameters["min_child_samples"] = 10
        locals().update(parameters)

```

```

In [44]: # custom function to run light-gbm model
def lgbm_rf(train_X, train_y, val_X, val_y, test_X):

    params = {
        "objective" : "regression",
        "metric" : "rmse",
        "num_leaves" : n_leaves,
        "learning_rate" : learn_rate,
        "bagging_fraction" : bagging_fraction,
        "feature_fraction" : feature_fraction,
        "bagging_freq":bagging_freq,
        'max_depth':-1,
        "min_child_samples" : min_child_samples,
        "boosting":"rf"
    }

    lgtrain = lgb.Dataset(train_X, label=train_y)
    lgval = lgb.Dataset(val_X, label=val_y)
    model = lgb.train(params, lgtrain, 3000, valid_sets=[lgval], verbose_eval=500,keep

    pred_test_y = model.predict(test_X, num_iteration=model.best_iteration)
    pred_val_y = model.predict(val_X, num_iteration=model.best_iteration)
    return pred_test_y, model, pred_val_y

In [45]: #!/time pred_test, model, pred_val = lgbm_rf(dev_X, dev_y, val_X, val_y, test_X)

In [46]: #!/final_score,sub_df = score()

In [47]: #!/plot_imp(model)

```

## 0.4 Iterative Testing

```

In [48]: for x in [x for x in range(4,10)]:
        x = 2**x
        print("testing for ",x)
        parameters["n_leaves"] = x
        locals().update(parameters)
        #!/time pred_test, model, pred_val = lgbm(dev_X, dev_y, val_X, val_y, test_X)
        final_score,sub_df = score()
        print("-----")

testing for 16
Training until validation scores don't improve for 300 rounds.
[300]      valid_0's rmse: 1.63636
[600]      valid_0's rmse: 1.63987
Early stopping, best iteration is:
[380]      valid_0's rmse: 1.62873
CPU times: user 6.27 s, sys: 255 ms, total: 6.53 s
Wall time: 1.93 s

```



1.6287219856869415

-----

testing for 32

Training until validation scores don't improve for 300 rounds.

[300] valid\_0's rmse: 1.63159

[600] valid\_0's rmse: 1.63221

Early stopping, best iteration is:

[400] valid\_0's rmse: 1.62434

CPU times: user 13.5 s, sys: 1.42 s, total: 15 s

Wall time: 5.16 s

1.624220219230191

-----

testing for 64

Training until validation scores don't improve for 300 rounds.

[300] valid\_0's rmse: 1.63217

[600] valid\_0's rmse: 1.63511

Early stopping, best iteration is:

[397] valid\_0's rmse: 1.62528

CPU times: user 15.3 s, sys: 471 ms, total: 15.8 s

Wall time: 4.34 s

1.625158693119286

-----

testing for 128

Training until validation scores don't improve for 300 rounds.

[300] valid\_0's rmse: 1.63217

[600] valid\_0's rmse: 1.63426

Early stopping, best iteration is:

[397] valid\_0's rmse: 1.62528

CPU times: user 21.8 s, sys: 187 ms, total: 22 s

Wall time: 5.59 s

1.6251567686026698

-----

testing for 256

Training until validation scores don't improve for 300 rounds.

[300] valid\_0's rmse: 1.63217

[600] valid\_0's rmse: 1.63426

Early stopping, best iteration is:

[397] valid\_0's rmse: 1.62528

CPU times: user 41.5 s, sys: 1.98 s, total: 43.5 s

Wall time: 13.1 s

1.6251567663041573

-----

testing for 512

Training until validation scores don't improve for 300 rounds.

[300] valid\_0's rmse: 1.63217

[600] valid\_0's rmse: 1.63426

Early stopping, best iteration is:

[397] valid\_0's rmse: 1.62528

```
CPU times: user 39.9 s, sys: 1.25 s, total: 41.1 s
Wall time: 11.4 s
1.625156766379118
-----
```

```
In [51]: for x in [2**x for x in range(4,10)]:
          print("testing for ",x)
          parameters["n_leaves"] = x
          locals().update(parameters)
          %time pred_test, model, pred_val = lgbm_rf(dev_X, dev_y, val_X, val_y, test_X)
          final_score,sub_df = score()
          print("-----")
```

```
testing for 16
[500]      valid_0's rmse: 1.73791
[1000]     valid_0's rmse: 1.74423
[1500]     valid_0's rmse: 1.74733
[2000]     valid_0's rmse: 1.75378
[2500]     valid_0's rmse: 1.75313
[3000]     valid_0's rmse: 1.75144
CPU times: user 26 s, sys: 1.01 s, total: 27 s
Wall time: 7.7 s
1.751441067818203
-----
```

```
testing for 32
[500]      valid_0's rmse: 1.81967
[1000]     valid_0's rmse: 1.83387
[1500]     valid_0's rmse: 1.8389
[2000]     valid_0's rmse: 1.847
[2500]     valid_0's rmse: 1.84379
[3000]     valid_0's rmse: 1.84142
CPU times: user 31.7 s, sys: 346 ms, total: 32.1 s
Wall time: 8.14 s
1.8414151694689804
-----
```

```
testing for 64
[500]      valid_0's rmse: 1.82008
[1000]     valid_0's rmse: 1.8346
[1500]     valid_0's rmse: 1.83987
[2000]     valid_0's rmse: 1.84797
[2500]     valid_0's rmse: 1.84474
[3000]     valid_0's rmse: 1.84241
CPU times: user 39.9 s, sys: 1.55 s, total: 41.5 s
Wall time: 11.5 s
1.842406748104294
-----
```

```
testing for 128
```

```

[500]          valid_0's rmse: 1.82008
[1000]         valid_0's rmse: 1.8346
[1500]         valid_0's rmse: 1.83987
[2000]         valid_0's rmse: 1.84797
[2500]         valid_0's rmse: 1.84474
[3000]         valid_0's rmse: 1.84241
CPU times: user 42.5 s, sys: 2.03 s, total: 44.5 s
Wall time: 12.8 s
1.842406748104294
-----

```

```

testing for 256
[500]          valid_0's rmse: 1.82008
[1000]         valid_0's rmse: 1.8346
[1500]         valid_0's rmse: 1.83987
[2000]         valid_0's rmse: 1.84797
[2500]         valid_0's rmse: 1.84474
[3000]         valid_0's rmse: 1.84241
CPU times: user 42.7 s, sys: 1.87 s, total: 44.6 s
Wall time: 12.9 s
1.842406748104294
-----

```

```

testing for 512
[500]          valid_0's rmse: 1.82008
[1000]         valid_0's rmse: 1.8346
[1500]         valid_0's rmse: 1.83987
[2000]         valid_0's rmse: 1.84797
[2500]         valid_0's rmse: 1.84474
[3000]         valid_0's rmse: 1.84241
CPU times: user 40.4 s, sys: 1.02 s, total: 41.4 s
Wall time: 10.9 s
1.842406748104294
-----

```

```

In [53]: for x in [x for x in range(16,33)]:
          print("testing for ",x)
          parameters["n_leaves"] = x
          locals().update(parameters)
          %time pred_test, model, pred_val = lgbm(dev_X, dev_y, val_X, val_y, test_X)
          final_score,sub_df = score()
          print("-----")

```

```

testing for 16
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63636
[600]          valid_0's rmse: 1.63987
Early stopping, best iteration is:
[380]          valid_0's rmse: 1.62873

```

```

CPU times: user 6.6 s, sys: 298 ms, total: 6.89 s
Wall time: 2.06 s
1.6287219856869415
-----
testing for 17
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63779
[600]          valid_0's rmse: 1.64154
Early stopping, best iteration is:
[399]          valid_0's rmse: 1.62945
CPU times: user 9.05 s, sys: 748 ms, total: 9.8 s
Wall time: 3.29 s
1.6294353614099684
-----
testing for 18
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63331
[600]          valid_0's rmse: 1.6355
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62482
CPU times: user 7.47 s, sys: 315 ms, total: 7.78 s
Wall time: 2.25 s
1.6248081114912876
-----
testing for 19
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63129
[600]          valid_0's rmse: 1.634
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62301
CPU times: user 7.62 s, sys: 250 ms, total: 7.87 s
Wall time: 2.17 s
1.6229723334486241
-----
testing for 20
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63121
[600]          valid_0's rmse: 1.63453
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62302
CPU times: user 8.19 s, sys: 370 ms, total: 8.56 s
Wall time: 2.48 s
1.622974978937571
-----
testing for 21
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63222
[600]          valid_0's rmse: 1.6333

```

```

Early stopping, best iteration is:
[401]          valid_0's rmse: 1.62429
CPU times: user 7.88 s, sys: 236 ms, total: 8.11 s
Wall time: 2.2 s
1.6242393262244306
-----
testing for 22
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63242
[600]          valid_0's rmse: 1.63471
Early stopping, best iteration is:
[403]          valid_0's rmse: 1.62476
CPU times: user 8.04 s, sys: 223 ms, total: 8.26 s
Wall time: 2.22 s
1.6246811229180314
-----
testing for 23
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63093
[600]          valid_0's rmse: 1.63254
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62409
CPU times: user 9.74 s, sys: 710 ms, total: 10.4 s
Wall time: 3.47 s
1.6239978885849027
-----
testing for 24
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63154
[600]          valid_0's rmse: 1.63518
Early stopping, best iteration is:
[405]          valid_0's rmse: 1.62478
CPU times: user 10.1 s, sys: 709 ms, total: 10.8 s
Wall time: 3.57 s
1.6246903455290411
-----
testing for 25
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63293
[600]          valid_0's rmse: 1.63101
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62506
CPU times: user 10.6 s, sys: 746 ms, total: 11.3 s
Wall time: 3.54 s
1.6249800387555666
-----
testing for 26
Training until validation scores don't improve for 300 rounds.

```

```

[300]          valid_0's rmse: 1.6318
[600]          valid_0's rmse: 1.62999
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62331
CPU times: user 10.4 s, sys: 695 ms, total: 11.1 s
Wall time: 3.49 s
1.6231936069407145
-----
testing for 27
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63265
[600]          valid_0's rmse: 1.63167
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62442
CPU times: user 8.72 s, sys: 247 ms, total: 8.97 s
Wall time: 2.43 s
1.6243252429876853
-----
testing for 28
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63268
[600]          valid_0's rmse: 1.63329
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62427
CPU times: user 11.2 s, sys: 673 ms, total: 11.9 s
Wall time: 3.49 s
1.6241598282203922
-----
testing for 29
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63265
[600]          valid_0's rmse: 1.63488
Early stopping, best iteration is:
[400]          valid_0's rmse: 1.62493
CPU times: user 10.1 s, sys: 454 ms, total: 10.6 s
Wall time: 3.03 s
1.6248069142765094
-----
testing for 30
Training until validation scores don't improve for 300 rounds.
[300]          valid_0's rmse: 1.63151
[600]          valid_0's rmse: 1.63329
Early stopping, best iteration is:
[397]          valid_0's rmse: 1.62449
CPU times: user 10.2 s, sys: 436 ms, total: 10.6 s
Wall time: 3.04 s
1.6243706858432774
-----

```

```

testing for 31
Training until validation scores don't improve for 300 rounds.
[300]      valid_0's rmse: 1.63242
[600]      valid_0's rmse: 1.6333
Early stopping, best iteration is:
[396]      valid_0's rmse: 1.62495
CPU times: user 10.3 s, sys: 400 ms, total: 10.7 s
Wall time: 3.03 s
1.624837110639381
-----
testing for 32
Training until validation scores don't improve for 300 rounds.
[300]      valid_0's rmse: 1.63159
[600]      valid_0's rmse: 1.63221
Early stopping, best iteration is:
[400]      valid_0's rmse: 1.62434
CPU times: user 9.41 s, sys: 146 ms, total: 9.56 s
Wall time: 2.45 s
1.624220219230191
-----

```